

Business proposal

A focused pilot for procedure readiness and targeted refreshers

The problem

A completed course does not establish whether an employee can apply every critical procedure step. When guidance changes, that completion record may also conceal outdated understanding. Supervisors need to know which steps still need help and which prior evidence remains applicable.

Our proposition

AgentX Learn connects approved guidance, learner diagnostics, targeted coaching and fresh scenario assessment. The agent selects an eligible activity using saved learner evidence. The backend grades the answer and checks readiness. When a procedure changes, a trainer reviews the update and approves which objectives require fresh practice.

First customer and use case

The proposed first buyer is an operations or finance lead in a small or medium-sized organization with recurring procedure questions and avoidable corrections. The initial use case is expense reimbursement. This customer hypothesis has not yet been validated through interviews or a paid pilot.

Stakeholder	Value to validate
Employee	Focused practice on an observed gap, with approved guidance close at hand.
Trainer / process owner	Evidence by objective, visible critical gaps and a saved human-review queue.
Operations lead	A traceable response to procedure changes without automatically repeating every objective.

Small initial adoption decision

Start with one process owner, one employee role and one procedure. Interview the owner and two or three employees about recurring errors and teaching effort. Proceed only if the problem is material and the owner can review the source, objectives and cases.

The working solution

Implemented prototype and observed fictional demonstration

The demonstration

The learner deliberately missed receipt evidence in a diagnostic. The live agent selected the approved receipt lesson, then a different scenario. Fresh cases established all four original objectives. A trainer then reviewed a revised policy permitting a written Finance exception when the supplier cannot replace a receipt. Three unchanged objectives carried forward; the receipt objective required an updated lesson and new case. Passing that case restored readiness for version two.

Why this is useful

Design choice	Practical distinction
Evidence, not completion alone	Readiness requires fresh assessment evidence and every critical objective.
Policy-version awareness	Historical results remain intact while affected objectives are reassessed.
Inspectable choices	Each new activity records selection mode, evidence at selection, source reference and outcome.
Human approval	The trainer reviews changed content and approves evidence carryover; the model cannot publish policy.

These are design distinctions, not claims of market exclusivity. No external competitor benchmark has been performed.

What is working today

The local prototype includes learner and trainer views, approved-source chat, persistent learning sessions, deterministic grading, a review queue, procedure revisions, targeted refreshers, a results dashboard and separate live/offline rehearsals. It runs with FastAPI, plain JavaScript and SQLite plus a configured model gateway.

Evidence achieved

The latest recorded checks passed 71 Python tests and 2 dashboard tests. Three consecutive gateway smoke runs passed after bounded JSON recovery. A browser rehearsal completed the original and updated fictional procedure. These results demonstrate implementation behavior; they do not establish improved retention, reduced trainer workload or production reliability.

Pilot and commercial path

Proposed evaluation and economics, not achieved business outcomes

A fair first pilot

Compare the workflow with the organization's existing training or document Q&A process using comparable unseen cases. Keep evaluation cases separate from training examples. Have a trainer review scoring and, where practical, score without knowing the condition. For a small volunteer group, report exploratory results and raw counts rather than broad causal claims.

Measure	Proposed decision criterion
Critical mistakes	Fewer missed mandatory steps on unseen cases; show per-person results.
Trainer effort	Lower active preparation, review and correction time without worse reviewed quality.
Grounded answers	Review supported and unsupported questions for factual accuracy, citations and appropriate referral.
Retention	Repeat equivalent unseen cases after a defined delay agreed with the process owner.
Operating cost	Track model calls, latency, hosting and content-maintenance effort per completed journey.

Commercial hypothesis

Offer a guided pilot, then consider an organization subscription with an onboarding/content-review service and usage limits. Pricing and willingness to pay remain unvalidated. Expansion to other procedures should follow measured benefit, not feature count.

Economic model

Estimate monthly capacity released from observed learner minutes saved and trainer hours saved, valued using the customer's own labor assumptions. Subtract model, hosting and continuing content-review costs. Three of four objectives carried in the demo does not mean 75% of time or cost was saved.

Risks and next milestone

The main risks are poor case design, unsupported explanations, weak transfer to real tasks and content-review overhead. Start with reviewed material and a narrow procedure. Managed identity, organization-specific access and public deployment are separate work. The next business milestone is a supervised real pilot.

Basis: supplied organizer submission screenshot; local README.md, PITCH.md, REHEARSAL.md and recorded test results. No market-size, customer-adoption or financial-return figures are asserted. Team code and repository URL are maintained in the email package pending confirmation.